

MANJARA CHARITABLE TRUST
RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI
 JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

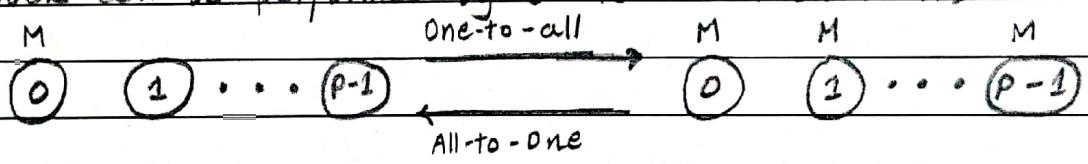
UNIT 3: PARALLEL COMMUNICATION

* One to All Broadcast :-

- Single processor sends identical data to all processors.
- Mostly used in parallel operations / algorithm.
- If 'p' is the number of processors and 'm' is the size of data then total 'mp' memory space will be used considering all processors.
- Each processor will have its own copy of data.

* All to one reduction :-

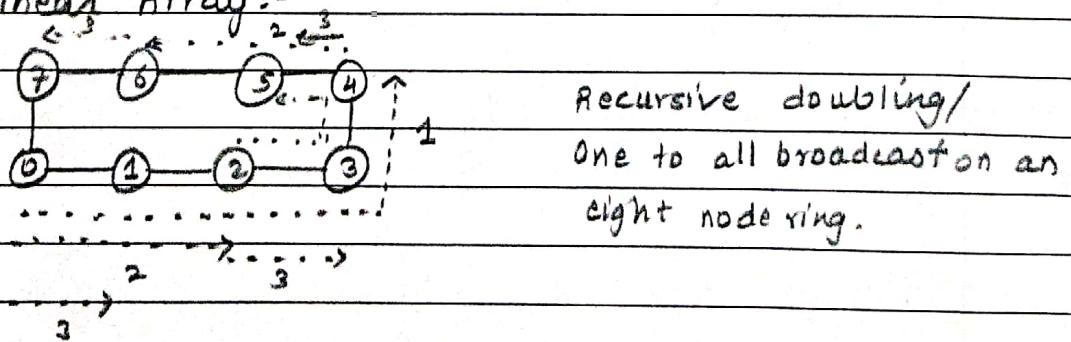
- Data from all processes are combined at single destination process.
- Operations like sum, maximum, minimum, product for a set of numbers can be performed by all to one reduction.



~ There are 3 topologies for the above 2 types of broadcasting :-

i) Ring/Linear Array (10). (ii) Mesh . (iii) Hypercube.

* Ring/ Linear Array :-



1. Sequential passing of a message from one process to another

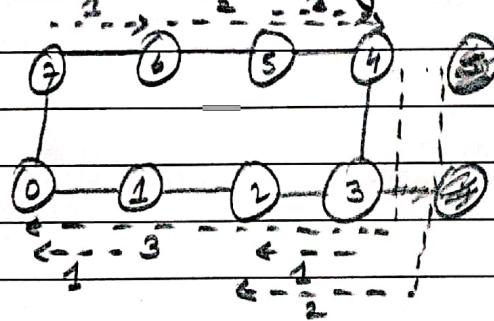
RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI.

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

causes congestion & is time consuming.

- 2. To overcome this problem in 1D array we use recursive doubling.
- 3. For sequential passing message passing takes $O(p-1)$ time whereas recursive doubling can do it in $O(\log p)$.
- 4. In recursive doubling given we start from 0^{th} node and we want to send message to all $p-1$ nodes then we first send message from 0 to $p/2^{\text{th}}$ node.
- 5. Now both 0^{th} and $p/2^{\text{th}}$ node will send message simultaneously in recursive manner.

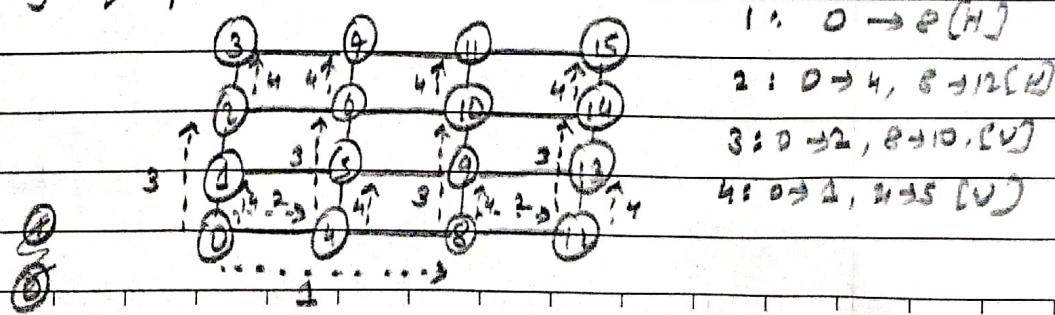
* Reduction on linear array:



- In reduction first all odd processors pass message to the next even processor;
- Then all even processor pass message to even processor.
- Finally $p/2^{\text{th}}$ node will pass message to destination node.

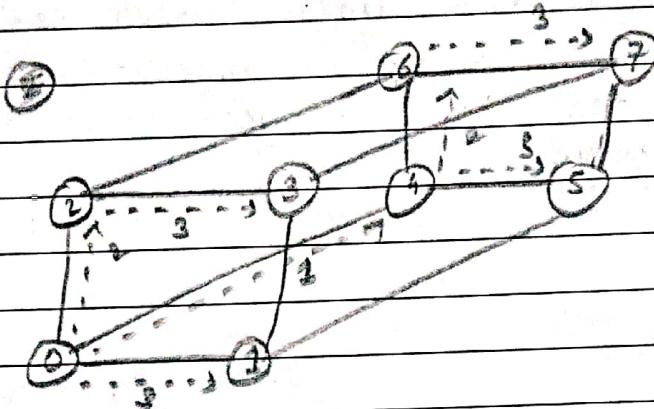
* Mesh:

- Each row & column of a square mesh p can be considered as linear array of \sqrt{p} nodes.



- Communication works in two phases: first row-wise and then column-wise.

* Hypercube :



- Similar to 2D mesh, 3D mesh works in two phases.
 - A hypercube with 2^d nodes can be considered a d-dimensional mesh.
 - We carry out d steps, one in each dimension.
 - communication starts from highest dimension bit.
- $\therefore 0 \rightarrow 4$ i.e. 000 to 100

* cost Analysis.

- For p processes, each message will be transferred at least $\log p$ times.
- Each transfer will have cost: $ts + tw \cdot m$
- Total time taken: $T = (ts + tw \cdot m) \cdot \log p$.

* All-to-All Broadcast & Reduction:-

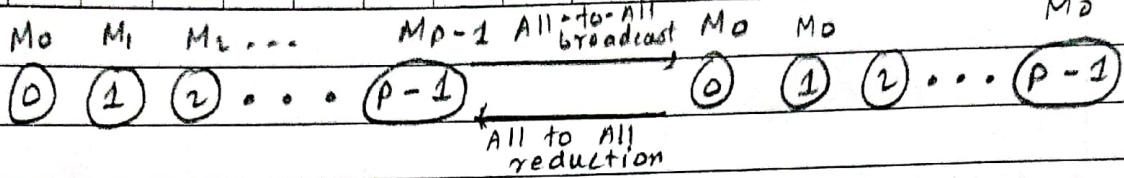
- In all-to-all all p nodes broadcast the message simultaneously.
- All processes may not contain the same message.
- In all-to-all reduction takes place on every node.

MGT Mp-1 Mp-1 Page No. Mp-1

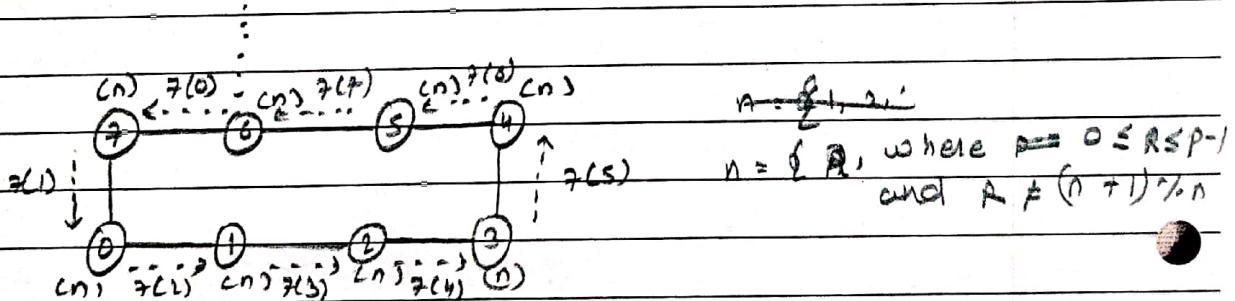
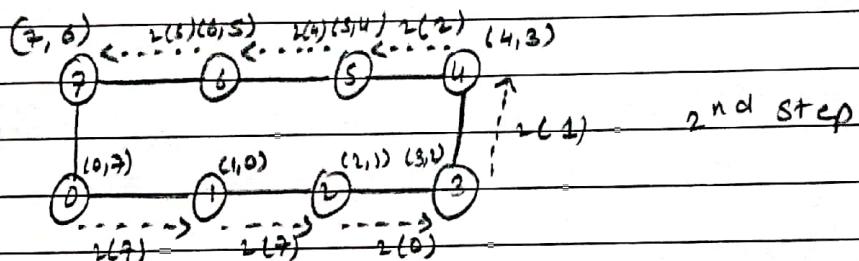
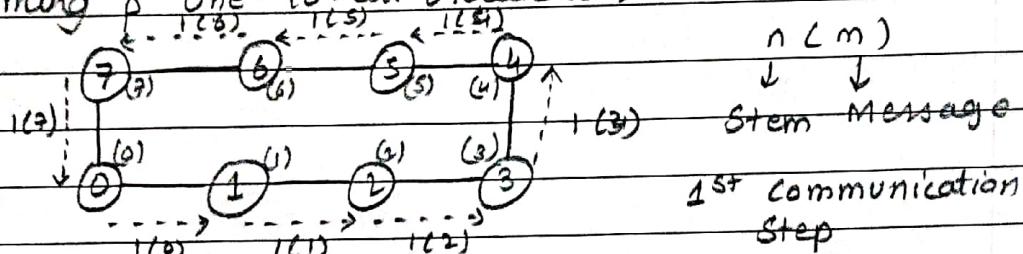
MANJARA CHARITABLE TRUST

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI.

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

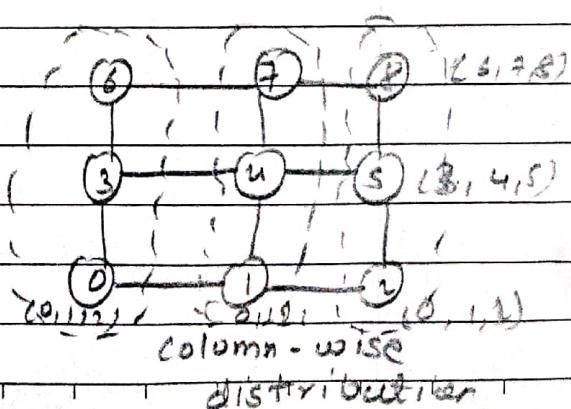
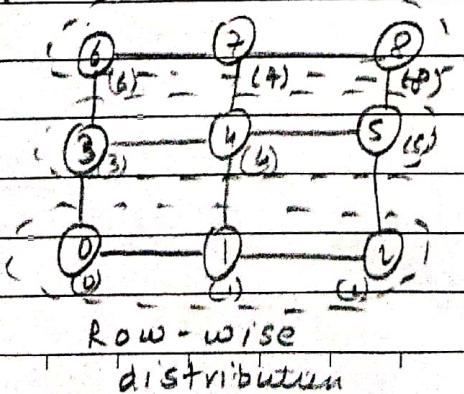


- One approach to perform A-2-A broadcast is by performing p one-to-all broadcast.



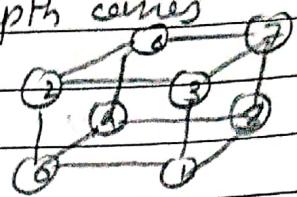
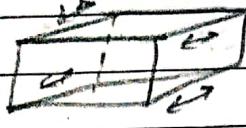
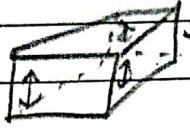
- For linear or ring it takes 7 rounds.

Mesh:



MANJARA CHARITABLE TRUST
RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI
JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- Hypercube:
- Same here just 3rd step which is depth comes



- * Cost Analysis :-

$$T = 2ts(\sqrt{p} - 1) + twm(p-1) \text{ for Mesh or Ring}$$

$$T = ts \log p + twm(p-1) \text{ for 3D.}$$

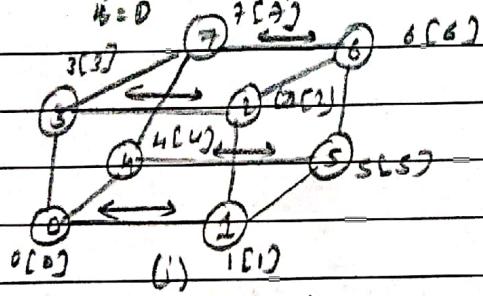
- All-reduce & Pre-fix sum operations :-

- All-reduce = All-to-one reduction + All-to-one broadcast.

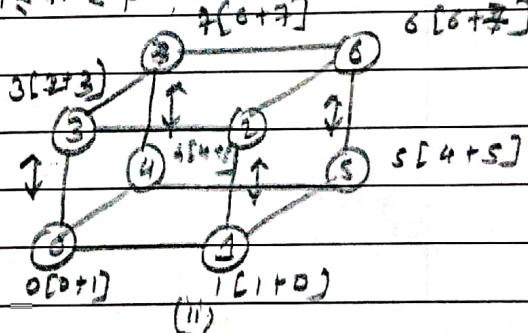
- After operation identical buffer size m are formed by combining mp buffers.

- Prefix sum :-

$$S = \sum_{i=0}^{k-1} n_i, 0 \leq i \leq k, 1 \leq n \leq p$$



Do same for
column
(iii)



Final prefix sum
(iv)

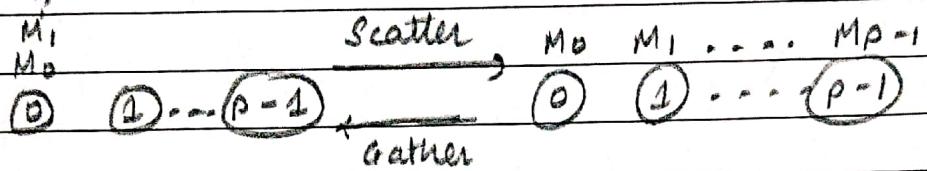
- Give example.

- Keywords :- Memory - Buffer, process - node,

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI.

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- * communication using MPI : Scatter & Gather
- one-to-all sends same data to all p nodes.
- In scatter, single node sends unique message of size 'm', to every other node.
- Scatter AKA one-to-all personalized communication.
- Gather operation is dual of scatter.
- Single node collects unique message from each node.



- Gather differs from all-to-one reduce as it does not reduce or combine collected data.

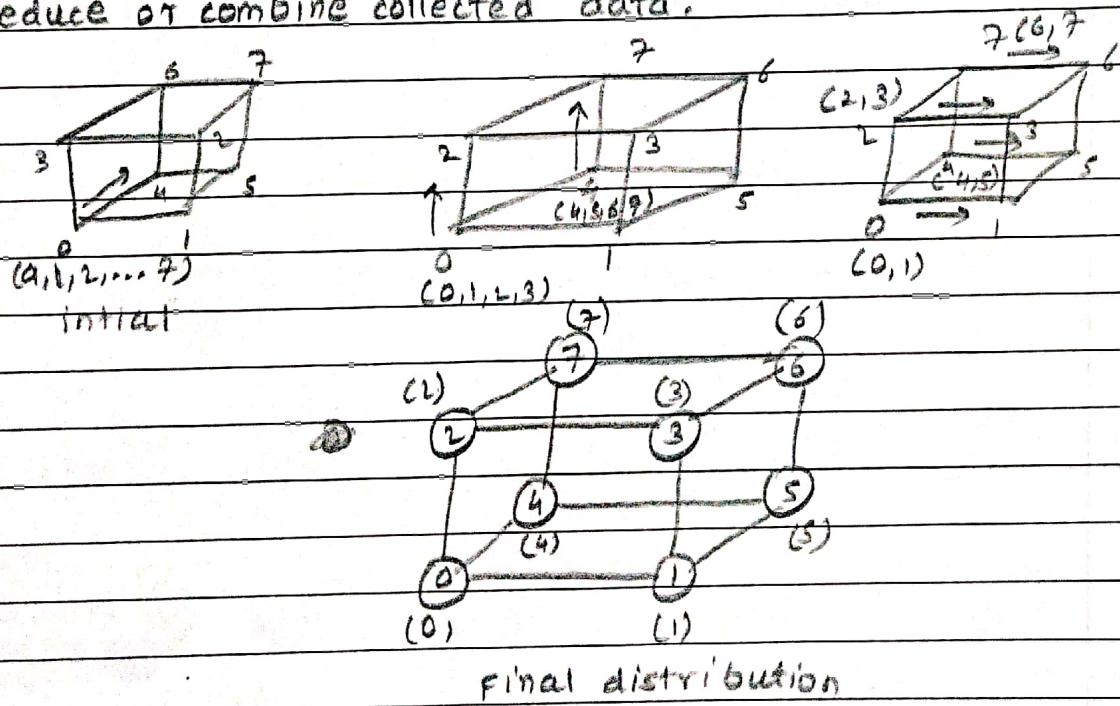


Fig : scatter op bration

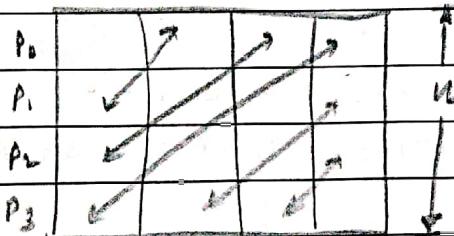
MANJARA CHARITABLE TRUST
RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI
 JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

Cost: $T = ts \log p + tw m(p-1)$

Same time for $2D \times 1$ lines.

* All-to-All personalized communication:

- Every node sends a distinct message of size m to every other node.
- | | | |
|---------------------------------|---------------------------------|---------------------------------|
| $M_{0,p-1}$ | M_{p-1,M_0} | M_{p-1,M_p} |
| $M_{0,1}, M_{0,0}$ | $M_{p-1,1}, M_{p-1,0}$ | $M_{1,0}, M_{0,0}$ |
| \circ \circ \dots $(p-1)$ | \circ \circ \dots $(p-1)$ | \circ \circ \dots $(p-1)$ |
- $\xrightarrow{\text{All-to-All personalized}}$



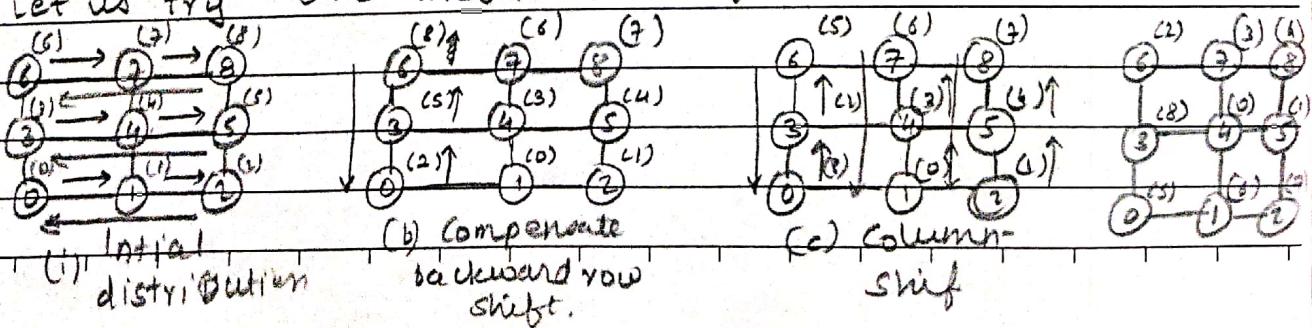
All-to-All personalized communication in a 4×4 matrix,

- It is similar to taking transpose of a matrix.

* Circular Shifts:-

- Used in matrix computation & string/image matching operations.
- It is a broader class of global communication model.
- In a circular q shift, node i sends data to $(i+q) \% p$ in a group of p nodes where $0 \leq q \leq p$.
- This is known as permutation.
- Algorithm :- 1) Circular shift row wise, 2) Circular shift only 1st row, 3) Circular shift all column. 4) Done.

Let us try 3×3 mesh with $q = 4$.



RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI.

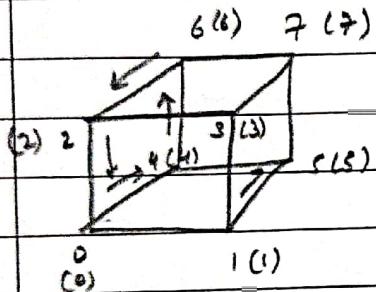
JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

$$\text{Cost: } T = (ts + mtw) (\sqrt{p} + 1)$$

$$i = (i + q_v) \cdot p$$

- Hypercube:

- For circular shift in hypercube values of i^{th} node are assigned to j^{th} node where j is binary reflected gray code.
- Two phases :- • First phase = 2×4 shift
 • Second phase = 1 shift.



UNIT 4: ANALYTICAL MODELING OF PARALLEL PROGRAMS

- Time Complexity of parallel algorithm depends upon :-
 i) Input Size. (ii) No. of processing units (iii) Communication speed of processors. (iv) Relative computation of processors.
 - * Sources of overhead in parallel programs :-
 - Doubling hardware resources may not double the performance.
 - Overhead typical execution involves Essential computation (serial), Excess computation (Parallel), Idling, interprocess communication :-
 1) Interprocess communication. \Rightarrow Idling (Load imbalance, Synchronization)
 2) Excess computation (Difference in computation by parallel program and serial program).
 - * Performance Measure & Analysis :-
 - Some metrics are :- 1) Efficiency. 2) Speedup. 3) Cost. 4) Parallel Overhead. 5) Execution Time.
 - Speedup :-
 - Captures relative benefit of solving a problem in parallel.
 - Given as :
$$S = \text{Speedup} = \frac{T_0}{T_p}$$

(Time taken by best sequential algorithm / Time taken by single processor)

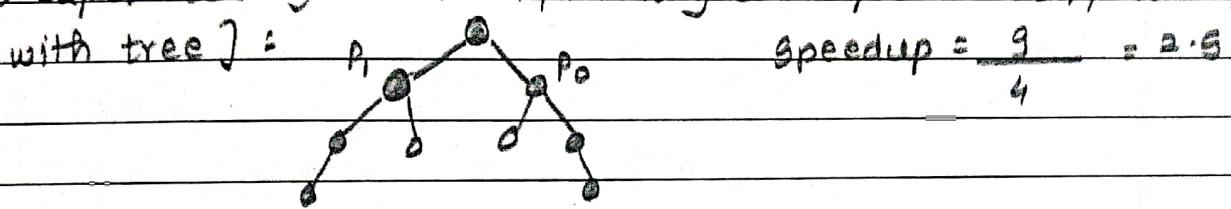
(Time taken by p processors)

 - Consider summation of n elements using n processing units.
 - Each processor is assigned a number to be added and in end a single processor will hold total sum.
 - Total addition will take place in $\log n$ steps.
 - Addition takes constant time as well interprocess communication.
- $\therefore T_p = \Theta(\log n) \quad S = \Theta\left[\frac{n}{\log n}\right]$
- Speedup.

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- Theoretically, speedup cannot be $> p$.
- Speed greater than p is possible only if each processing element spends time $< T_S/p$.
- In practice speedup $\geq p$ is observed, this is superlinearity effect.
- Happens when serial does more work than parallel.
- Two reasons:- 1) Superlinearity due to caches [Fast Access],
2) Superlinearity due to exploratory decomposition [If working with tree] :-



$$\text{Speedup} = \frac{9}{4} = 2.25$$

- * Efficiency :- Ratio of Speedup to no. of processing elements.
- $S/E = \frac{S}{P}$ Ideally: $S = p$ and $E = 1$.
In practice $S < p$ and $E < 1$.

It is the measure of time fraction of time for which a processor is fully employed.

- Cost :- Product of parallel runtime & no. of processors used.
- Cost given as: $W = T_p * p$.
- A parallel system is cost optimal if it has same asymptotic growth as the best known sequential algorithm.
- Execution Time :-
- Time elapsed from the moment parallel computation starts to the moment it ends.
- Serial runtime is T_S and parallel is T_p .

* Amdahl's law:-

- Consider a fixed computational workload W and a number of processors p .
- Let execution rate of i processors be R_i so for

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

sequential processors ($i=1$) $R_1 = 1$ & $R_n = n$

- Assume sequential workload is αw ($0 \leq \alpha \leq 1$) and parallel workload is $(1-\alpha)w$.
 $\therefore w_1 = \alpha w$ & $w_n = (1-\alpha)w$.

- Total Time execution is :

$$T_n = \frac{w_1}{R_1} + \frac{w_n}{R_n} = \frac{\alpha w}{1} + \frac{(1-\alpha)w}{n}$$

$$T_n = \alpha w + T_n = w \left[\alpha + \frac{1-\alpha}{n} \right] = w \left[\frac{\alpha(n-1)+1}{n} \right]$$

- Speed up is :-

$$S_n = \frac{T_1}{T_n} = \frac{w/1}{w \left[\frac{\alpha(n-1)+1}{n} \right]} = \frac{n}{1+(n-1)\alpha}$$

This is Amdahl's law, when $n \rightarrow \infty$

$$S_{\infty} = \frac{1}{\alpha}$$

Even when $n \rightarrow \infty$, Speedup cannot be increased beyond $(1/\alpha)$ this is a bottleneck for multiprocessors.

- * Gustafson's law :-

- Consider workload is increased to n -node machine.
 $\therefore w' = \alpha w + (1-\alpha)n w$

- Speed up is :-

S' = Single processor execution
 Parallel process execution

$$S' = \frac{(\alpha w + (1-\alpha)n w)/1}{\frac{\alpha w}{n} + \frac{(1-\alpha)n w}{\alpha}} \dots \left[T = \frac{w}{P} \text{ for } i=1 \right] \quad \left[T = \frac{w}{2} = w \right]$$

$$S' = \frac{\alpha + n - n\alpha}{\frac{\alpha}{n} + (1-\alpha)} = \frac{(\alpha + n - n\alpha)/n}{(\alpha/n + (1-\alpha))} = \frac{\alpha + n - n\alpha}{\alpha + n - n\alpha}$$

$$S' = \frac{\alpha + (1-\alpha)n}{\frac{\alpha}{n} + (1-\alpha)} = \frac{\alpha n + (1-\alpha)n^2}{\alpha + n(1-\alpha)}$$

MANJARA CHARITABLE TRUST

Page No. _____

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

It simplifies as $S'n = \alpha + (1-\alpha)n$

- So Gustafson's law says that true parallel power is only achievable when large parallel problem is applied.

* Effect of Granularity on performance.

- No. of task divisions \propto Granularity.
- Fine granularity :- Decomposition into large no. of tasks.
- Coarse granularity :- —————— || — small — || — —
- Degree of concurrency \propto Granularity.
- In practice we do fine granularity.
- Using less processing elements than available is called scaling down of a parallel system.
- If there are n inputs and p processing elements and $p < n$ then we can design a parallel algo that works for n processing elements by using n/p virtual processing elements.
- As processing element decrease by n/p the workload at each processing unit increases by a factor of n/p .
- Granularity impacts :-
 - 1) It affects performance of parallel computers.
 - 2) Finer granularity increases Speedup.
 - 3) However synchronization overload negatively affects performance.

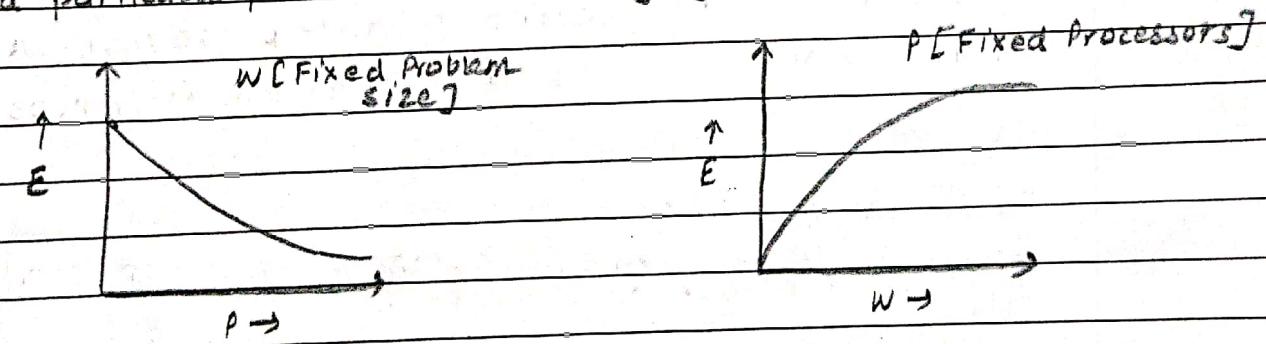
* Scalability of parallel systems :-

- Capacity to increase speedup in proportion to number of processing elements.
- Reflects parallel system's ability to utilize increasing processing resources effectively.
- Efficiency is given as :-

$$E = \frac{S}{P} = \frac{T_S}{P T_P} = \frac{1}{1 + \frac{T_P}{T_S}}$$

~~NOT~~
MANJARA CHARITABLE TRUST
RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI
 JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- Total overhead T_o is an increasing f^n of p .
- Every parallel program always has a sequential part that runs for time t_s .
- \therefore Total overhead = $(p-1) \times t_s$
- Hence Overhead increases linearly with p , however, it may increase superlinearly due to communication overhead, idle time, excess computation.
- \therefore As $p \uparrow E \downarrow$ goes down $E \propto 1/p$.
- Scale more then efficiency shall decrease.
- Efficiency may increase if problem size is increased keeping number of processors constant.
- Isoefficiency metric of scalability:
- Helps determine best algorithm/architecture combination for a particular problem without trying them all.



- Isoefficiency function helps determine ease with which a parallel system can maintain constant efficiency.
- A parallel system is cost-optimal iff $p T_p = \Theta(W)$ where W is overhead overhead.
- And $W = \Omega(T_o(p+t_p))$
- A parallel system is cost-optimal if its overhead function does not exceed its problem size asymptotically.

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- $w(p)$ is asymptotic lower bound on isoefficiency.
- isoefficiency f^n is optimal only if degree of concurrency is $\Theta(w)$.

* Minimum execution time & Minimum cost :-

- As the number of processors increase the parallel runtime continues to decrease and asymptotically approaches a minimum value, or goes on increasing after a min val.
- We give min execution times as:-

$$\frac{d}{dp} T_p^{\min} = 0$$

This gives the no. of processing elements for which T_p is min.

- However max. no. of processing elements is bounded by degree of concurrency.

$$T_p^{\min} = \frac{w + T_0(w, c[w])}{c[w]}$$

* Asymptotic Analysis of parallel programs:-

Example of adding n nos. with n processing elements

$$T_p = \Theta(\log n) \quad S = \Theta(n/\log n)$$

* Matrix Computation:-

Matrix - vector Multiplication :-

procedure MAT_VEC(A, X, Y)

begin

for i := 0 to n-1 do

 y[i] := 0;

 for j := 0 to n-1 do

 y[i] := y[i] + A[i][j] * X[j];

 endfor

endfor

end MAT_VEC

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI
JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- Sequential algo needs n^2 multiplications & additions.
 - There are 3 ways of doing this parallelly.
 - i) Row-wise 1D Partitioning :-

Row-wise ID Partitioning :-

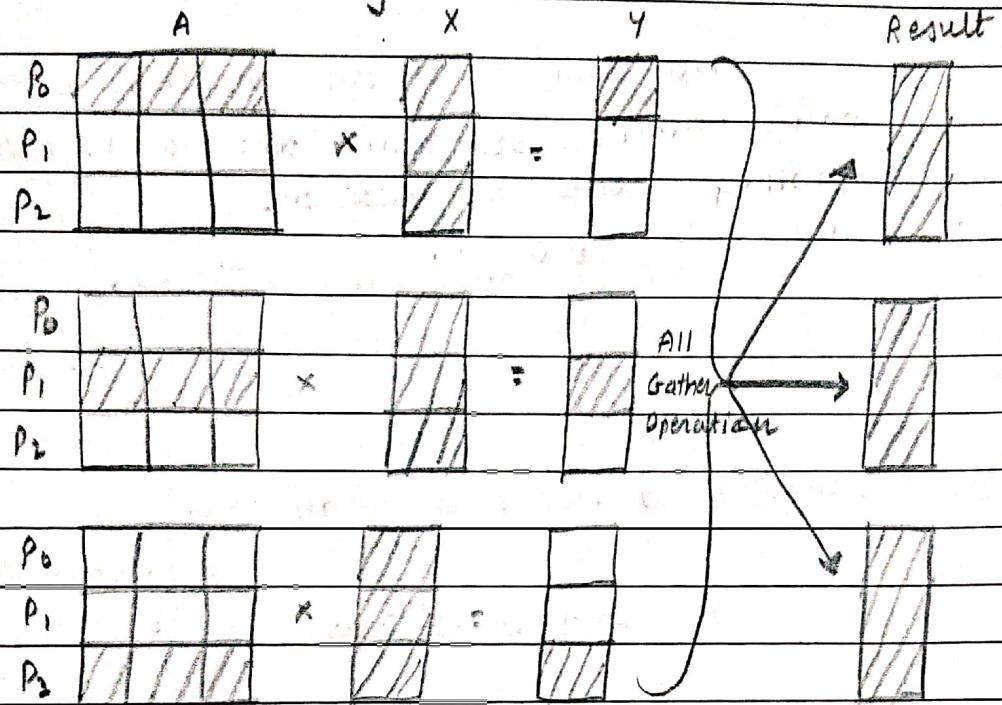
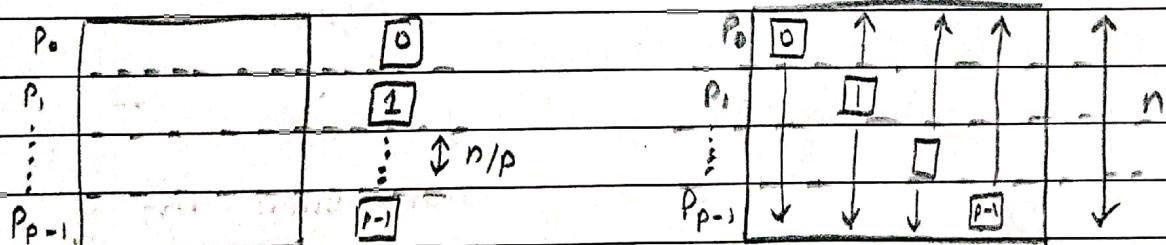
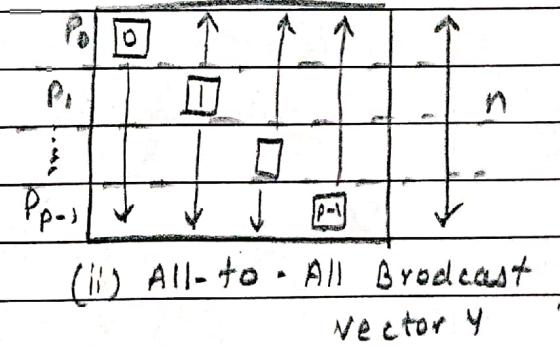


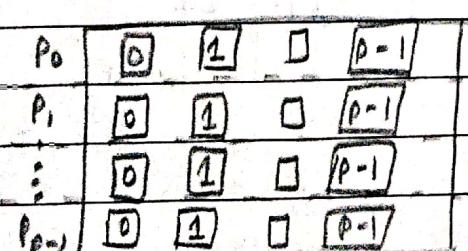
Fig: Computation Scheme for Rowwise Data decomp.
Matrix A Vector X



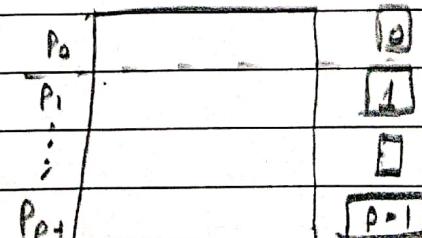
(1) Initial partition



(ii) All-to-All Broadcast



(iii) Entire vector distributed



(iv) Final Results.

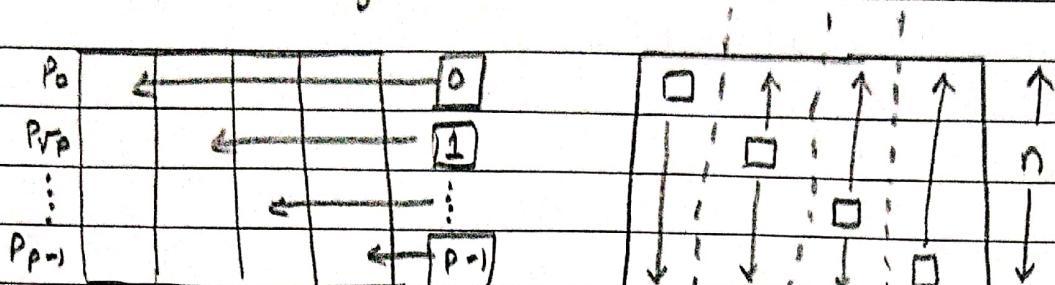
It is $O(n^2)$ so cost optimized

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

2) Columnwise 1 - D.

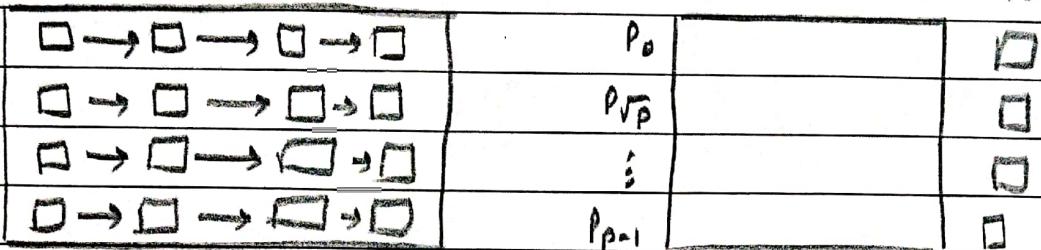
3) 2D - Partitioning :-



(i) Distribute along diagonal

(ii) One-to-all broadcast

Matrix A Vector y



(iii) All-to-one reduction

(iv) Final distribution,

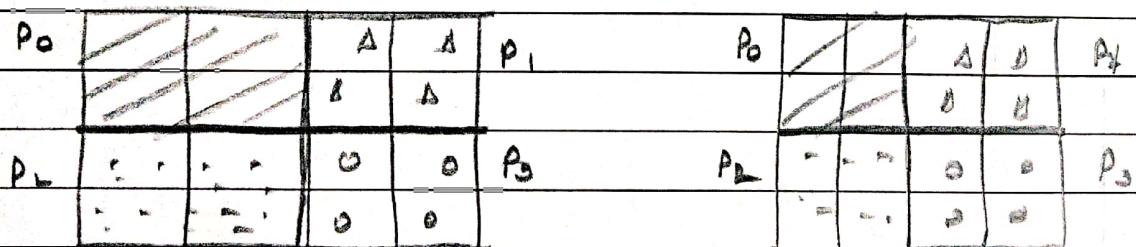
cost is not optimized. It is $O(n^2 \log n)$

* Matrix - Matrix Multiplication :-

- Sequential algo is same with only addition of 3rd loop from 0 to $n-1$.

Matrix A

Matrix B

Divide in blocks assign to P_s

$$C_0 = P_0 \times P_{0,B}, \quad C_1 = P_1 \times P_{1,B}$$

we use left shift & Up shift on processes.

Explain with examples.

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

UNIT 5: CUDA Architecture

* Introduction to GPU :-

- It is the age of parallel computing as many users in today's world do multi-tasking on their phones, tablets, laptops etc.
- To provide novel, fast and rich-experience to the user, developers need to understand & use various parallel platforms.
- CPU performance increasing number of cores/transistors & other features and always has been a reliable way of enhancing performance but not the only way.
- GPUs:-
- Specialized electronic circuit. Helps alter memory rapidly to accelerate creation of images that is given as output to the display device.
- GPUs are in :- Mobile Phones [Adreno/Mali], PCs [AMD/NVIDIA], Consoles [PS4/S, XBOX].
- GPUs are great at image processing & controlling computer graphics.
- They are highly parallel.

• CUDA:-

- Compute Unified Device Architecture.
- Launched by NVIDIA in 2007 as an programming interface providing parallel computation using GPUs.

• CUDA Architecture:-

- Before CUDA, vertex & pixel shaders were used for parallel computing.
- Pixel shader is a GPU component that can operate per pixel.

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- Vertex shader is also like a GPU component like pixel but it is assembly language specific, used for geometrical operations.
- CUDA was to be used for general-purpose computing, rather than being specifically for graphics.
- 3 Device-level APIs: i) OpenCL ii) DirectX iii) CUDA Driver API
- Language Level: 3 :- i) Fortan, ii) C++, iii) C.
- DirectX runs standalone, Rest three need CUDA Driver.
- Then we go software-level i.e CUDA support in OS Kernel.
- Then we go hardware-level i.e CUDA parallel compute engine inside NVIDIA GPUs.

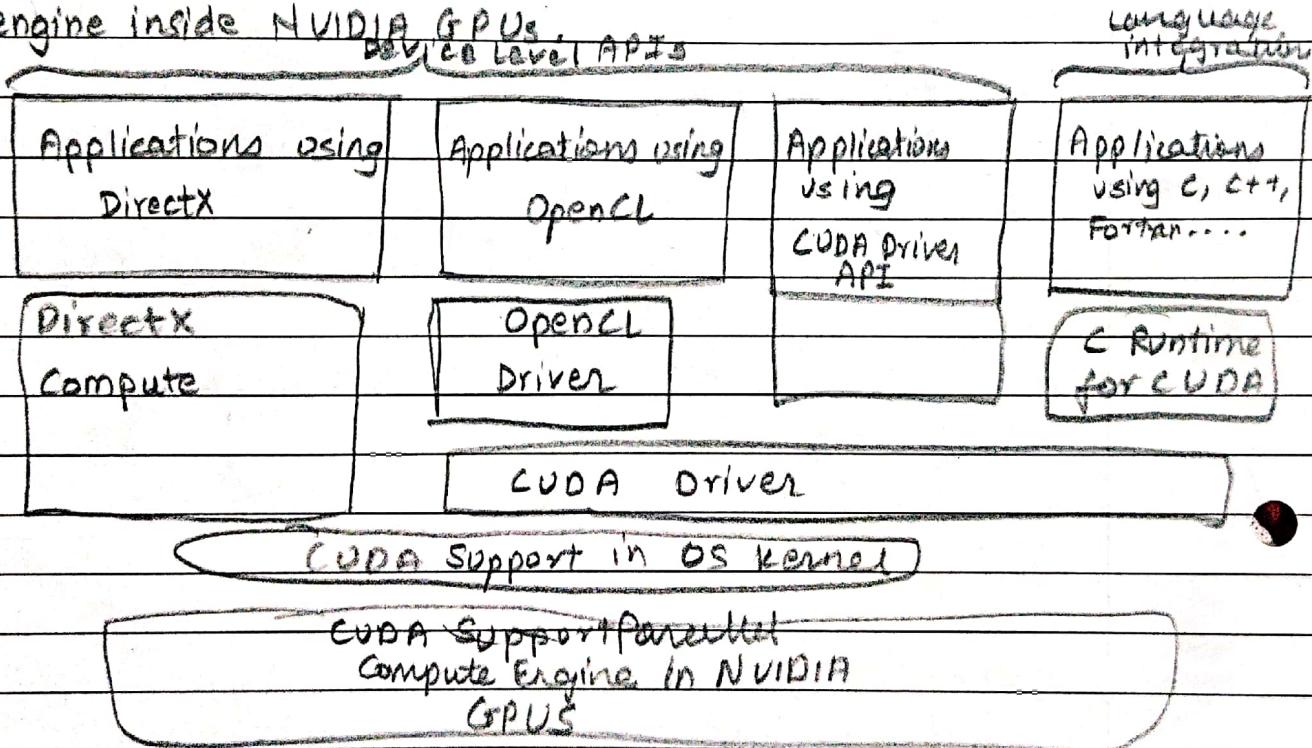
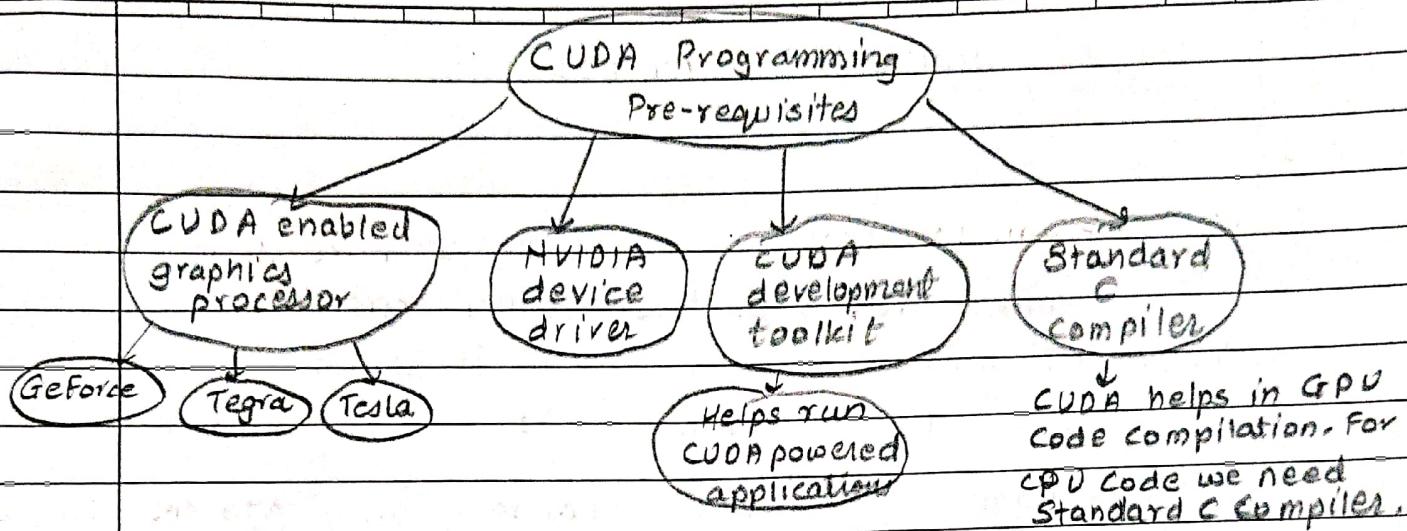
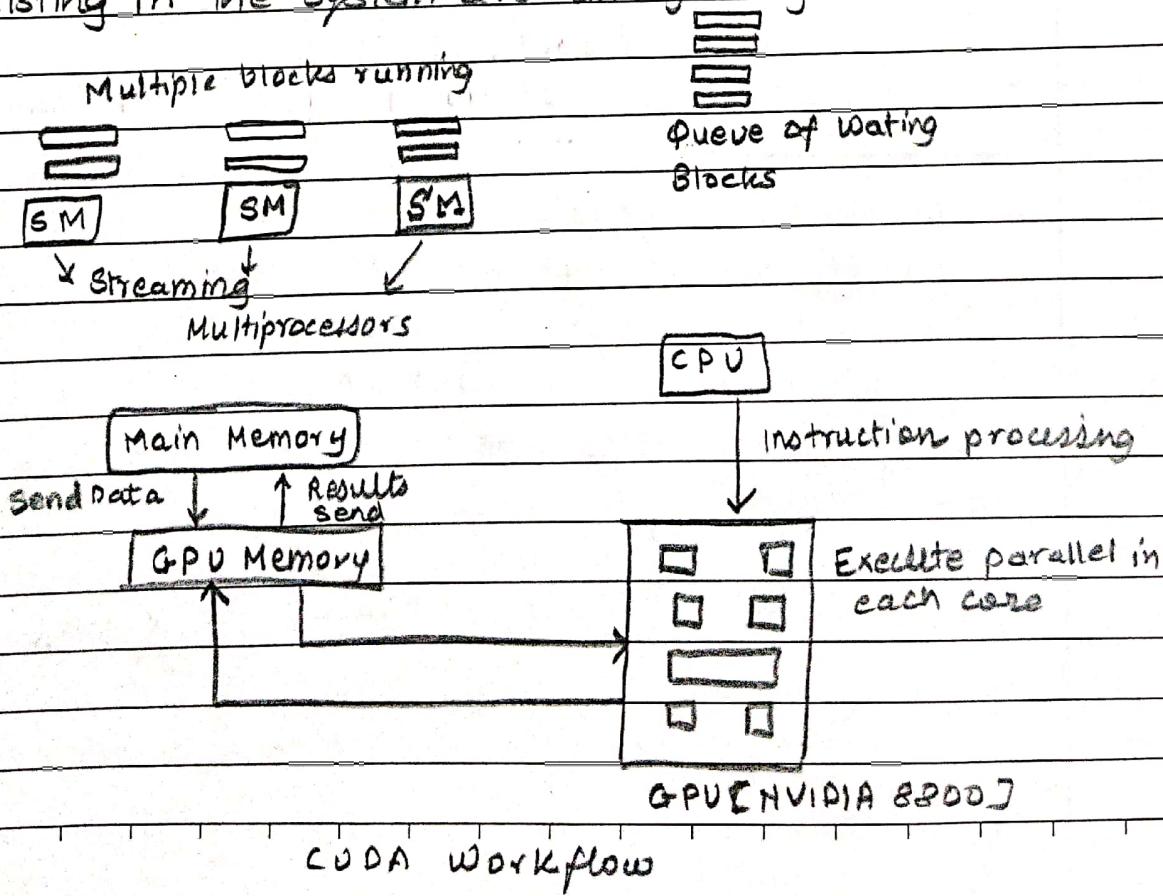


Fig: CUDA Architecture.

- Execution units on GPUs are allowed read & write access on a cache known as shared memory.
- CUDA driver consists of parallel computing kernels and functions.
- CUDA C++



- Working of CUDA C/C++ Program:-
- Serial code executes in a host [CPU] thread.
- Parallel code executes in many concurrent device [GPU] threads across multiple parallel processors called as streaming Multiprocessors.
- Hence to perform parallel execution of code, at least one CPU is the minimum requirement and one or more GPU existing in the system are arranged by CUDA architecture.

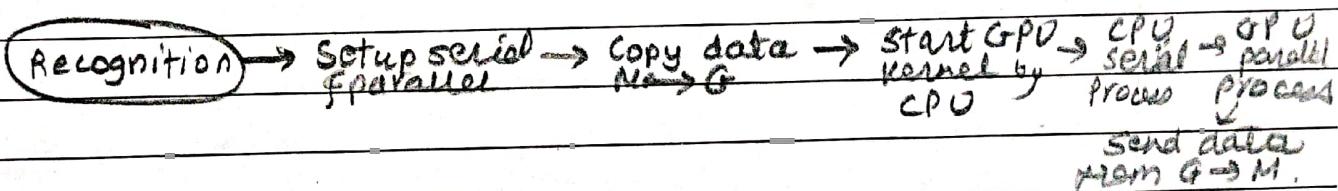


RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

Flow of CUDA C code :-

1. CPU recognizes serial code & parallel code.
2. CPU threads initialized to run serial code & GPU threads for parallel.
3. GPU threads copies data from main mem. to GPU mem for parallel processing.
4. CPU initiates GPU compute kernel.
5. CPU does serial code.
6. GPU's CUDA cores execute kernel in parallel.
7. GPU thread copy data result to main memory.



* Example of CUDA Kernel programs:-

```

#include <iostream>
__global__ void kernel(void) {
}
int main() {
    kernel<< 1, 1>>x);
    printf("Hello World!\n");
    return 0;
}
    
```

- Global keyword helps call kernel function using host.
- Global helps run the kernel fn on GPU
- we use NVCC compiler to compile this code instead of gcc.
- `kernel<< 1, 1>>x()` is a CUDA specific syntax that specifies a call to device code.
- Parameters inside "`<< >>`" are called execution configuration

MANJARA CHARITABLE TRUST
RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI
 JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- Execution configuration determines no. of threads & blocks.

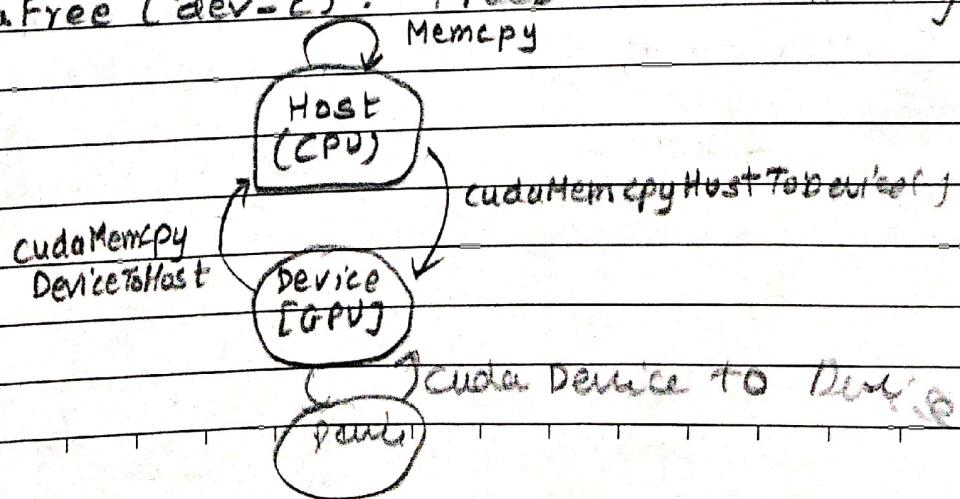
- * CUDA Memory Model :-

```
#include <iostream>
global __ void (int* a, int b, int*c) {
    *c = a + b;
}

int main(void) {
    int * c;
    int * dev_c;
    cudaMalloc((void**) dev_c, sizeof(int));
    add << 1, 1 >> (2, 7, dev_c);
    cudaMemcpy(&c, dev_c, sizeof(int), cudaMemcpyDeviceToHost);
    printf("2 + 7 = %.d \n", c);
    cudaFree(dev_c);
}
```

- The main memory methods used :-

- 1) `cudaMalloc (void** dev_c, sizeof(int))` :-
 Allocates memory dynamically for on device for result.
- 2) We allocate it in terms of bytes using `sizeof()`.
- 3) `cudaMemcpy (&c, dev_c, sizeof(int), cudaMemcpyDeviceToHost)`,
 copy dev_c to c.
- 4) `cudaFree (dev_c)` :- Free allocated memory for dev_c.



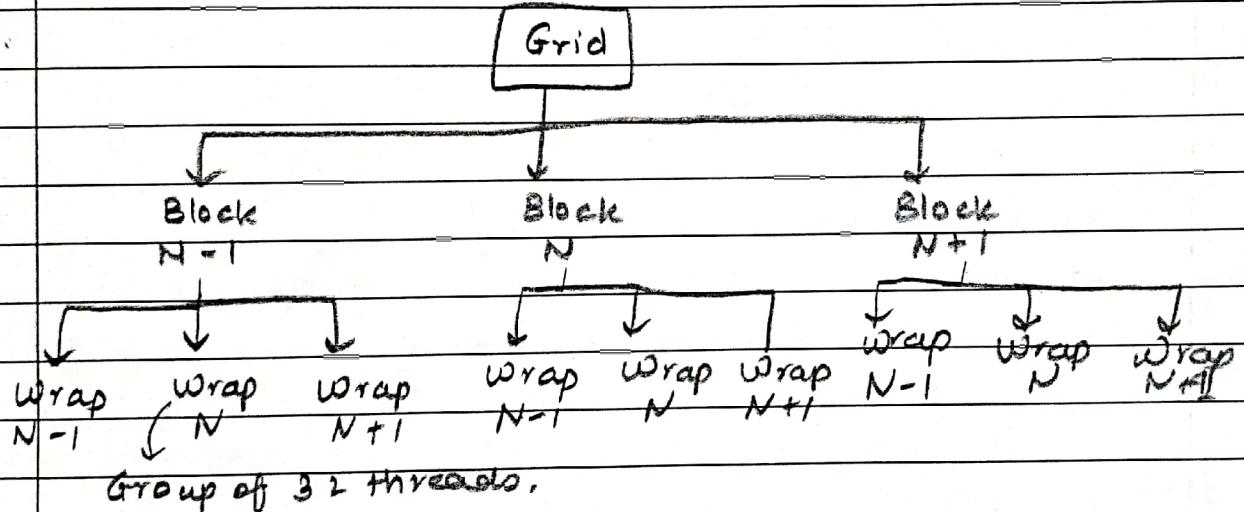
RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- Surround every cuda call with HANDLE_ERROR().
- It is a utility macro that detects errors,
- It helps in debugging.

* Manage Communication & Synchronization:-

- CUDA is ideal when little to no inter-thread/block communication is needed.
- For inter-thread on chip resources with explicit primitives are used.
- For inter-block multiple kernels are invoked, using off-chip global memory.
- Methods Used : `cudaDeviceSynchronize()`
for Synchronization `cudaDeviceReset()`



- Simultaneous Multiprocessors switch between wraps.
- Blocks run in any order, only a subset of blocks executes anytime.
- Blocks run on SMs, and are allocated in a round-robin format.
- Wraps exhibit SIMD.

- Parallel Programming in CUDA C [Theory] :-

I] Threads :-

- Fundamental unit of a program.
- Thread is a lightweight process, handles 1 task at a time.
- Helps bring concurrency in application.
- Give example of vector addition using threads.
- Thread info stored in `threadIdx` i.e a built-in CUDA variable.
- `ThreadIdx` is accessed as :: `int tid = threadIdx.x`
- Thread id is unique within each block.
- Each thread that executes kernel can be referenced by its ID.

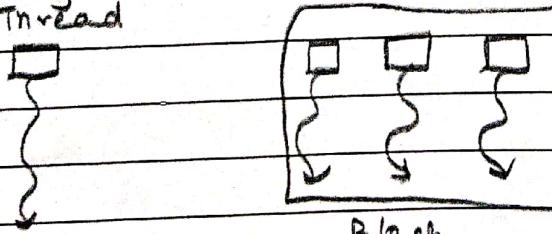
II] Blocks :-

- Threads are organized into blocks, blocks are referred as `blockIdx` in CUDA runtime.
- Threads within block synchronize, on a modern GPU a block may contain 1024 threads.
- Using block ID & thread ID we can make unique thread ID for each kernel.
- So kernel $\ll<$ Number.of.blocks, Number.of_THREADS $\gg>$ \gg \ll is how we call a kernel.
- \therefore Kernel $\ll<$ 256, 1 $\gg>>$ gives 256 blocks.
- Specifically $\ll<$ Number.of.blocks, Number.of.Threads $\gg>>$ \ll per grid \gg per block

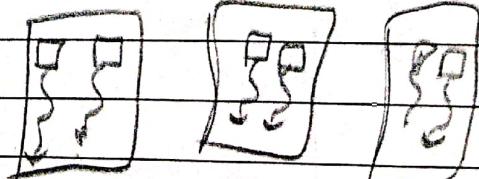
III] Grid :-

- Grid allows to define a block in two dimensions.
- Mostly used in matrix multiplication / other operations.

Thread



Block



Grid

UNIT 6: HIGH PERFORMANCE APPLICATIONS

- * Parallel Sorting :-

- * Sequential algorithm :-

```
procedure BUBBLE-SORT(n)
```

```
begin
```

```
for i := 0 to n-1 do
```

```
    for j := i+1 to n-i-1 do
```

```
        compare-exchange (aj-1, aj);
```

```
    endfor
```

```
endfor
```

```
end BUBBLE-SORT()
```

- * Odd-Even Transposition :-

- This algorithm sorts an array of n elements in n phases (if n is even), each phase requiring $n/2$ comparisons.

- It alternates between odd & even phase, in odd phase, only elements at odd indexes are compared, and in even; even.

```
procedure ODD-EVEN(n):
```

```
begin
```

```
    for i := 0 to n-1 do
```

```
        begin if i is odd then
```

```
            for j := 0 to n/2-1 do
```

```
                compare-exchange (aj, aj+2)
```

```
            if i is even then
```

```
                for j := 1 to n/2-1 do
```

```
                    compare-exchange (aj, aj+2)
```

```
        endfor
```

```
    end ODD-EVEN
```

[only write
end for
statements
that have
begin]

- The algorithm can be made parallel.

```
procedure ODD-EVEN-PAR(n)
```

```
begin
```

```
for id := process-label
```

```
for i:=1 to n do
```

```
begin
```

```
if i is odd then
```

```
    if id is odd then
```

```
        compare-exchange_min(id+1);
```

```
    else
```

```
        compare-exchange_max(id-1);
```

```
    if i is even then
```

```
        if id is even then
```

```
            compare-exchange_min(id+1);
```

```
        else
```

```
            compare-exchange_max(id-1);
```

```
    endfor
```

```
end ODD-EVEN-PAR
```

- Here $p=n$ and complexity is $O(n^2)$ which is not cost optimal.
- If $p=\log n$ then optimal cost performance is achieved.

* Distributed computing:-

- Document classification:-

- Process of assigning predefined labels/categories to documents which include unstructured & semi-structured.

- Procedure:-

- 1) set of pre-classified documents used as training set.

- 2) A classification scheme is derived.

- 3) Scheme is refined using testing process.

MANJARA CHARITABLE TRUST

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- * New docs are now classified,
- Doc classification on sequential CPUs takes lot of time.
- GPU can be used as parallel unit for this task.
- Parallel CNN in CUDA / MPI can be used here.

* High Performance Frameworks:-

- Kubernetes :-
 - An open source system for handling development lifecycle.
 - Helps in automating deployment, scaling and management of containerize application that also support distributed computing.
 - Helps in not compromising security or reliability.
 - Saves time on infrastructure management.
 - Acts as an orchestrator for containers.
- Parallel computing for AI / ML - Kubeflow:
 - ML Toolkit for Kubernetes.
 - Keeps workflows simple, portable & scalable.
 - Supports Jupyter notebooks & tensorflow model convinient.
 - Adds portability in ML Deployments.

* Scope of parallel Computing :-

- 1) Engg Engineering & Design :- Complex Math solving, Optimization.
- 2) Scientific Application :- DNA structures, Astrophysics, Bioinformatics.
- 3) Commercial :- Stock, Data Mining, Large scale servers.
- 4) Computer system :- Intrusion detection, Modern automobile, P2P Network.

* Parallel Merge Sort :

- Divide array into two parts. 1st part = array elements with odd index & 2nd = array elements with even index.
- Consider another array B.
- We define merge(A, B) as:

$$\text{merge}(A, B) = \{a_1, b_1, a_2, b_2, \dots, a_n, b_n\}$$

If $A = \{1, 2, 3, 4\}$; $B = \{5, 6, 7, 8\}$ then

$$\text{Merge}(A, B) = \{1, 5, 2, 6, 3, 7, 4, 8\}$$

- Now, $\text{Join}(A, B) = (\text{Merge}(A, B), \text{Odd-Even}(A, B))$

- Algorithm:

>>> procedure Odd Even(A, B, S)

begin

if A and B are of length 1 then

Merge A and B using compare-exchange operation

else

begin

compute Sodd and Seven in parallel do

Sodd = Merge(Aodd, Bodd)

Seven = Merge(Aeven, Beven)

Seven-odd = Join(Sodd, Seven)

end

end Odd Even

- Example:-

- Let $S = \{2, 3, 6, 10, 15, 4, 5, 8\}$; $A = \{2, 6, 15, 4\}$

$B = \{3, 10, 5, 8\}$;

$A_{\text{odd}} = \{2, 15\}$; $A_{\text{even}} = \{6, 4\}$; $B_{\text{odd}} = \{3, 5\}$; $B_{\text{even}} = \{10, 8\}$

- Now, $\text{Merge}(A_{\text{odd}}, B_{\text{odd}}) = \{2, 3, 5, 15\} \{2, 3, 5, 15\}$

$\text{Merge}(A_{\text{even}}, B_{\text{even}}) = \{6, 10, 4, 8\}$

$\text{Merge}(A, B) = \{2, 6, 3, 10, 5, 8, 15, 4\}$

Now, $\text{Join}(A, B) = [\text{Merge}(A, B), \text{Odd-Even}(A, B)]$.

We already know Merge (A, B), Join Odd-Even will sort (A, B) and give the final ans:-

$$\text{odd-even } (A, B) = \{2, 3, 4, 5, 6, 8, 10, 15\}$$

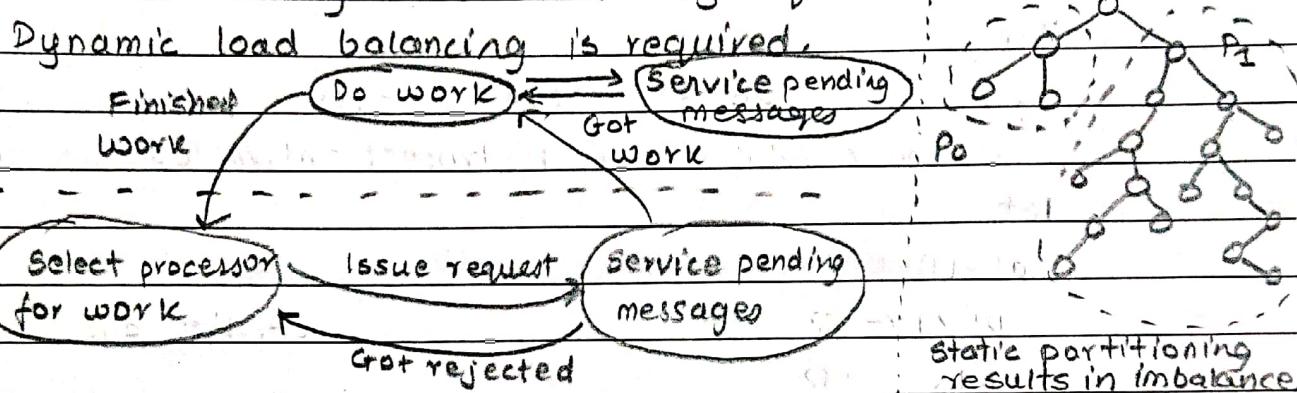
- Parallel merge-sort takes $\Theta(\log n)$ time on using n processor, hence total time: $n * \Theta(\log n) = \Theta(n \log n)$
- It is cost-optimal.

* Parallel DFS :-

- Idea :- Different sub-tree can be searched concurrently.
- It is necessary to estimate size of subtrees.
- Dynamic load balancing is required.

Processor Active

Processor Idle



- Work is split by splitting stack in two.
- Neither of split stack should be small, but both equal.
- Node splitting :- Select nodes from bottom of stack.
- Stack splitting :- Select nodes from each level of stack.
- Stack splitting is favored.
- Several ways to do load balancing :-
- 1) Asynchronous round robin :- Counter for each processor, requests made in round-robin style. $V(p) = O(p^2)$ [Poor]
- 2) Global round robin :- Global counter + round robin. $V(p) = O(p)$ [Poor performance due to limited resources; contention]
- 3) Random polling :- Randomly select processor. $V(p)$ is not bounded [Desirable performance]

RAJIV GANDHI INSTITUTE OF TECHNOLOGY, MUMBAI,

JUHU VERSOVA LINK ROAD, VERSOVA, ANDHERI (WEST), MUMBAI - 53.

- P_0 has all work, and has a weight associated with it.
- It will partition its work and give it to another processor.
- In doing so it loses half of its weight, and weight is gained by recipient processor.
- After completion of work, P_0 will return to its normal weight & termination is signalled.
- Work W is split into ψw & $(1-\psi)w$ and a constant $\alpha (0 < \alpha \leq 0.5)$ is used such that $\psi w > \alpha w$ & $(1-\psi)w > \alpha w$,
- α sets lower bound for load imbalance.

* Parallel BFS ::

- Key ideas :- 1) Do BFS Level-By-Level not vertex by vertex. [Level Synchronization]. 2) Process entire level in parallel.
 - Let us see the pseudocode :-
- ```

ParallelBFSGE(V, E, S, V)
 D[V] ← ∞
 D[S] ← 0
 L ← 0 // Level Counter
 F0 ← {S} // Add source to frontier
 while Fl ≠ ∅ do
 Fl+1 ← ∅ // Next level frontier
 processLevel(G, Fl, Fl+1, D)
 L ← L + 1 // Go to next level
 return D

```

- source: J. Bailey*
- ```

processLevel(G, Fl, Fl+1, D)
    if |Fl| > q0 then // If bag is big then split
        (A, B) ← bagSplit(Fl)
        spawn processLevel(G, A, Fl+1, D)
        processLevel(G, B, Fl+1, D)
    else // If bag ain't big
        sync
        for v ∈ Fl do
            par-for (v, w) ∈ E do
                if D[w] = ∞ then
                    D[w] ← l + 1
                    bagInsert(Fl+1, w)
    
```
- Frontier holds all the nodes of a single level.
 - While $F_l \neq \emptyset$: the span is not more than diameter of graph.
 - We use divide-and-conquer while splitting the bag that holds frontier.
 - A bag is a multi-set data structure.
 - Cost: $W = O(|V| + |E|)$ Total cost = $O(d \times \log |V \times E|)$

diameter